

BAYESIAN PHYLOLINGUISTICS

Simon J. Greenhill^{1,2}, Paul Heggarty¹ and Russell D. Gray^{1,2,3}

- ¹. Department of Linguistic and Cultural Evolution, Max Planck Institute for the Science of Human History, 07745 Jena, German.
- ². ARC Centre of Excellence for the Dynamics of Language, Australian National University, Canberra, ACT 0200, Australia.
- ³. School of Psychology, University of Auckland, Auckland 1142, New Zealand.

1. INTRODUCTION

Change is coming to historical linguistics. Big, or at least “big-*ish*” data (Gray and Watts 2017), are now becoming increasingly available in the form of large web-accessible lexical, typological and phonological databases (e.g. Greenhill et al. 2008, Bowerman 2016, Moran 2014, Haspelmath 2014, Bickel et al. 2017) and the soon to be released Lexibank, Grambank, Parabank and Numeralbank (<http://www.shh.mpg.de/180672/glottobank>). This deluge of data is way beyond the ability of any one person to process accurately in their head. The deluge will thus inevitably drive the demand for appropriate computational tools to process and analyze the vast wealth of freely available linguistic information.

Historical linguistics has long dabbled with computational and quantitative approaches (e.g. Chrétien 1943). The search for a rigorous methodology to make accurate inferences about subgrouping and timing of language splits first led to a quantitative subgrouping approach called lexicostatistics and extended those to dating with an approach called glottochronology. However, severe shortcomings were quickly identified with both lexicostatistics and glottochronology and led to an almost puritanical rejection of these approaches. More recently, new Bayesian phylogenetic methods from evolutionary biology – which do not share the fatal shortcomings of lexicostatistics and glottochronology – have been applied to linguistic questions.

In this chapter we review the history of quantitative approaches to language subgrouping and dating, and then turn specifically to Bayesian phylogenetic methods and their utility for historical linguistics. We focus on four main questions: what is Bayesian phylolinguistics, why does this approach typically focus on lexical data, how is it able to estimate divergence dates, and how reliable are the results?

2. EARLY QUANTITATIVE APPROACHES

Morris Swadesh (Swadesh 1950, 1952, 1955) proposed a simple computational approach to building language trees, which he called *lexicostatistics*. He argued that the number of cognate words shared between two languages in basic vocabulary was a good indicator of degree of relationship. The logic is simple: successively grouping languages by amount of cognates they shared would give rise to a family tree with more similar languages grouped together. Swadesh further realized that the degree of similarity shared between two languages was also an approximate measure of how much time had elapsed since they had separated. Making a direct analogy to radioactive decay, Swadesh argued that languages lost similarity at a relatively constant rate over time. These rates were naturally only applicable to a standardized sample of vocabulary that they were calculated from, so Swadesh proposed a series of standardized 100- and 200- item wordlists. For example, if two

languages shared C cognates and there was a constant “clock” rate r , then the age of divergence t could be estimated by:

$$t = \log C / 2 \log r$$

Swadesh (1950) initially suggested that the clock rate r was 85% every 1000 years based on the differences between old and modern English. This figure was later revised to 81% every 1000 years based on a survey of rates in the basic vocabulary of historically attested languages (Lees 1953). For example, if two languages shared 81% cognates on a 200-item word list, then the estimated age would be 1000 years, while two languages sharing 66% might be expected to have diverged 2000 years ago. This approach to dating the divergence of related languages became known as *glottochronology*.

Glottochronology was first applied by Swadesh to the Salishan languages (Swadesh 1950). However, the promise of using linguistics to date prehistoric population expansions was so tempting that glottochronology was rapidly adopted by linguists and used on languages around the world (see Hymes 1960 for a review). The findings were widely read outside linguistics with eminent anthropologists carefully discussing the methodology (Kroeber 1955) and arguing that it revolutionized prehistory (Murdock 1964).

However, critics were quick to note some serious methodological flaws, with the key flaw being that language change is not clock-like. Languages can vary substantially in their rates of lexical replacement. For example, a critical study showed that glottochronology would estimate that Old Norse and Icelandic diverged less than 200 years ago. This age is far younger than the historically attested age of 1000 years (Bergsland and Vogt 1962). In fact, rather than the retention rate of 81%/1000 years proposed by glottochronology, these rates varied by around 15-20% (Bergsland and Vogt 1962). The finding of substantial rate variation in languages was so damaging to glottochronology that historical linguistics largely rejected quantitative methods. Today, lexicostatistics and glottochronology are seen as textbook examples of bad linguistic methodology, and we are told that “linguists do not do dates” (McMahon and McMahon 2006).

In biology, the scenario started in a similar vein but played out rather differently. Biologists had developed a method to build trees from pairwise similarity (e.g. Sokal and Michener 1958). Divergence times could then be estimated from these trees using a “molecular clock” rate (Zuckerkandl and Pauling 1965); for example, Kimura (1968) estimated a rate of one base pair change every 1.8 years. Just as in linguistics, biologists were quick to notice the problem of rate variation which could cause incorrect estimates of both the relationships on the family tree (Felsenstein 1978), and the estimated ages (Kirsch 1969).

Unlike historical linguists, who largely rejected computational approaches, biologists developed methods for handling variation in rates. One method, known as *nonparametric rate smoothing* (Sanderson 1997), takes the observed rates of change in the data, and ‘smoothes’ these to fit around calibration information. Another method – currently the best – is known (Drummond et al. 2006) as the *relaxed clock* (in contrast to what might be called the *strict clock* assumption of glottochronology). In the relaxed clock model, rates for each branch are drawn from a log-normal or exponential distribution with parameters describing the magnitude and variation in rates estimated from the data. This procedure allows each branch to have its own rate and for the data to inform the analysis about the magnitude of rate variation across the languages.

3. WHAT IS BAYESIAN PHYLOGENETICS?

There is a proud tradition of building language family trees in historical linguistics that was popularized by August Schleicher in the nineteenth century (Atkinson and Gray 2005), but dates back to at least the seventeenth century (List et al. 2016). These language family trees are typically constructed in a qualitative manner from innovations in lexical, phonological, and morphological characters. However, despite the fact that implicit in the comparative method is some kind of optimization procedure, traditional historical linguists do not use an explicit optimality criterion to select the best tree, nor do they use an efficient computer algorithm to search for the best tree. This is surprising given that the task of finding the best set of trees is inherently a combinatorial optimization problem of considerable computational difficulty. In abstract mathematical terms, at least, for just five languages there are in theory 105 ways of subgrouping them in a rooted bifurcating tree. For 10 languages this number grows to over 34 million possibilities (Felsenstein 1979). Therefore, even between just the ten or so major branches of Indo-European, for instance, the theoretical number of possible tree configurations is enormous.

Traditional language family trees suffer from two additional limitations: their branching patterns reflect only a relative chronology and they contain no information about how strong the support is for any proposed subgroupings. Bayesian phylogenetic methods provide a useful *supplement* to the comparative method. They enable us to build trees with explicit estimates both of branch lengths and of subgrouping uncertainty in an objective, repeatable manner. These trees can then be used to evaluate subgrouping hypotheses (e.g. Gray et al. 2009, Greenhill and Gray 2012), date language divergences (e.g. Gray et al. 2009, Bouckaert et al. 2012), estimate ancestral states (e.g. Maurits & Griffiths 2014, Haynie & Bown 2016), infer rates of change in lexical and grammatical traits (e.g. Pagel et al. 2007, Dediu 2011, Greenhill et al. 2017), test hypotheses of functional dependencies in linguistic features (e.g. Dunn et al. 2011), and infer geographic homelands and migration routes (e.g. Bouckaert et al. 2012, Grollemund et al. 2016, Lee & Hasegawa 2013). So what exactly are Bayesian phylogenetic methods?

Bayesian Phylogenetic methods use Bayes' Theorem to make probabilistic inferences about phylogenetic trees and their model parameters. The approach infers a sample of trees that have a high probability of explaining the patterns in the data under a specified model of character change. In technical terms, the approach calculates a posterior probability distribution of trees $P(\text{tree}|\text{data})$ as a function of the prior probability of a tree $P(\text{tree})$ and the likelihood of the data (data) given the model of character change. In the analysis, the model specifies the number of parameters to be estimated for the rate/s of character replacement (whether these characters are cognate lexical items, or DNA nucleotides, or other data). The models can range from simple models with just one rate to more complex models where different types of character replacements have different rates. The model can also include complexities such as rate variation between different characters by fitting them to a distribution of rates (e.g. a gamma distribution). When applied to Swadesh lists, for instance, rates of cognate replacement can differ for different meanings in the list: across Indo-European, replacement rates are far slower in the meaning 'name' than for 'dirty', for example.

To estimate the posterior probability distribution from the data and the model, Markov Chain Monte Carlo (MCMC) methods are typically used. In this procedure, the search through tree and parameter space starts with a random tree and arbitrary values for the model parameters and branch lengths. A likelihood score is calculated for this tree and set of parameters. Then a new tree and model parameters are proposed. If the likelihood score is better, this new proposal is accepted. If the likelihood is worse, then the proposal will be accepted with a probability determined by the ratio of the likelihood of the proposed tree divided by the existing one (i.e.

really bad trees are unlikely to be kept). This process is continued for a very large number of steps (generations). The MCMC chain explores tree and parameter space incrementally, finding the set of trees that best fit the data given the model. The tree topology, branch lengths and model parameters are saved at intervals so that the progress of the chain can be assessed.

When the analysis starts the initial estimates of topology, branches and parameters are essentially random guesses and are therefore poor inferences of the true parameters. This initial phase is called ‘burn-in’ and is commonly removed from the analysis later. By repeatedly proposing and keeping ever better trees and parameter estimates the MCMC chain should end up converging on a region where the most probable trees are sampled in proportion to their posterior probability.

In practice, it is sensible to use multiple runs with different random starting trees to check that the analysis has converged on this region. For a recent review of Bayesian phylogenetics in biology that includes a good discussion of some of the practical issues with MCMC searches see Nascimento et al. (2017), while Greenhill and Gray (2009) provide a detailed overview in a linguistics context.

A key feature of Bayesian phylogenetic inference is that the result of the analysis is not a single tree but rather a set of trees and their model parameters, sampled in proportion to their posterior probability. This set of trees is often summarized in a consensus tree. Consensus trees are normally depicted with the posterior probability of each clade (subgroup) shown on the branch. This number is the percentage of trees in the posterior distribution that contains that branch. Thus we end up not just with a single optimal tree, but rather with a set of trees that lets us evaluate the extent to which the data supports a particular clade (given the model and the prior).

Figure 1 shows how this approach can be adapted for linguistic inferences. First, some data must be selected. To date this has mainly been cognacy in basic vocabulary. The rationale for this choice is discussed in the next section. At this point we simply note that in principle other kinds of linguistic data can be used in phylogenetic analyses, and have been (see Greenhill et al. 2010, Greenhill et al. 2017), and that combined analyses are also possible. With lexical data, the next and in many ways most crucial and time-consuming step involves coding the data for cognacy. Here the comparative method is king. We want real cognates, not mere “look-alikes”. While considerable progress has been made on automating cognate coding – up to 89% accuracy – these procedures are best viewed as assisting rather than replacing the linguist (List et al. 2017). Once the lexical data have been coded for cognacy, the cognate sets can be converted into a matrix suitable for phylogenetic analyses (Step 2). The matrix might be either a multistate or a binary one, depending on the model of character change that is going to be selected. In the multistate matrix, each semantic slot is a character (column) and the value in each cell reflects whichever cognate class that language uses in that meaning. In the binary matrix, each cognate set has its own column and the value in the cell reflects the presence (1) or absence (0) of that cognate in that meaning slot in that language (see figure 1).

Step 3 involves the specification of the prior. This can include information on the likely distribution of trees, model parameters, branch lengths and the like. One advantage of the Bayesian approach is that information from other sources can be included in the prior to aid the best inference. For example, the age of ancient texts or the timing of known historical events can be used to calibrate the trees to help date divergence times (see section 4). Inferences from other kinds of data (e.g. phonological innovations) could also be included in the prior to aid subgrouping inferences. Step 4 entails selecting a model of character change. Models can range from the simple single rate model

with strict clock to models that allow rates to vary (Yang 1993) and a relaxed clock. There is a tradeoff in model selection. If the model is too simple, the inferences may be inaccurate. If a model is too complex (“over-parameterised”), then the excess of parameters to be estimated will inflate the uncertainty in estimating each parameter and thus limit the power of the analysis (Burnham and Anderson 1998). In careful Bayesian analyses the performance of different models should be evaluated using Bayes factors to determine which model performs best. In practice we have found that the covarion model (Penny et al. 2001), where sites can switch between fast and slow rates on different parts of the tree, often outperforms simpler models such as the Continuous-Time Markov Chain model (which only allows one rate of cognate gain and loss) or the stochastic Dollo model (where cognates gain only be gained once). Similarly – and this will be of no surprise to critics of glottochronology (see below, §3.5) – the relaxed clock model outperforms the strict clock (see below, §4, on these “clock” models).

Step 5 is the MCMC search through tree and parameter space. Because each step in the Markov chain is necessarily strongly correlated with the previous step, samples from the chain need to be taken many generations apart (e.g. 1,000 or even 10,000 generations). The chain also needs to be run long enough to get well past the initial “burn-in” phase (where parameter estimates have yet to converge, see figure 2), and multiple runs with different random seeds are needed to check for convergence.

The final step in the Bayesian analysis is to summarize the tree topology and the key parameters of interest, such as divergence dates. While consensus trees are often used to summarize the inferences about the tree topology, in our experience linguists often ignore the posterior probabilities on the branches and treat all branches in the consensus tree equally. This is a serious mistake. The real result of the analysis is not a single tree but the posterior set of trees. A key feature of Bayesian phylolinguistics is that this set reveals the strength of support in the data for various subgroupings (given the prior and the model). A useful way to visualize this uncertainty is to plot all the trees from the posterior distribution on top of each other in a “densitree” (Bouckaert 2010). Figure 3 shows a densitree for Central Pacific languages. The tree was constructed from cognate-coded basic vocabulary from the Austronesian Basic Vocabulary Database (Greenhill et al 2008), using a binary covarion model. The densitree shows that some subgroups such as Eastern Polynesian are reliably recovered. It also reveals considerable uncertainty in the placement of many languages, perhaps as a consequence of the conflicting signal caused by dialect networks (Gray et al. 2010).

4. WHY USE “THE LEXICON”?

Historical linguists do not generally see the lexicon as their data-set of choice. The lexicon is the level of language most exposed to borrowing, and it is also not immune to chance resemblances that may appear to go back to a common original form, but do not (e.g. Spanish *mucho* and English *much*). The comparative method insists instead on other types of language data, deemed much more conclusive of language relationships: form-to-meaning correspondences across extensive morphological paradigms (e.g. the parallels first identified between the case endings in Sanskrit, Greek and Latin); and repeated, consistent sound correspondences, linked by regular and natural ‘laws’ of sound change between them. So why have most Bayesian phylogenetic analyses nonetheless drawn their comparative language data from the lexicon? And can the results be trusted? To answer these questions we first need to clarify what exactly is meant by ‘lexical data’. We also need to break this down into three separate issues about the suitability of the lexicon for different purposes: (1) to establish *whether* languages are related within a family; (2) to recover the *phylogeny* of an already established family; and (3) to research a family’s *chronology*.

4.1 The Lexicon for Establishing Relatedness?

Swadesh (1950, 1952, 1955 – see §2 above) first drew up his famous lists to target meanings that he assumed were the most stable and resistant to borrowing. In the early heyday of what Swadesh called lexicostatistics, this methodology morphed into an optimistic assumption that even a minimal level of apparent matches between the Swadesh lists for two languages could be presumed to reflect deep shared cognates rather than loanwords, and thereby demonstrated a deep ‘stock’-level relationship (see Heggarty 2010: 307-309). Even the core lexicon is by no means entirely borrowing-free, however, so claims to establish relatedness from lexical lists alone have never convinced. Basic vocabulary is still used in exploratory studies, as a fertile hunting ground where deep cognates are most likely to survive. But that has to be followed up by the orthodox comparative method, to either confirm or dismiss whether any apparent matches reflect a relationship of common descent (cognacy), or not. Likewise, a Bayesian phylogeny or ‘family tree’ is meaningful only for languages that do actually form a family with each other, and only when based on data that are probative of phylogenetic relationships. Such data lie not in lexis *per se*, but in cognate status, which can only be established by the comparative method. So Bayesian phylogenetic methods are not proposed here as a method to try to establish whether languages are related in the first place. On the contrary, they rely on the comparative method to establish that in advance, so as to create the input data they need: cognacy judgements. In any case, a simple lexicostatistical *count* of putative ‘matches’ has nothing to do with Bayesian phylogenetic analyses and their explicit *modelling* of descent with modification through time.

4.2 Why Use ‘The Lexicon’ for Phylogeny?

Once the comparative method has already established a family, one can then move on to the second task: recovering the structure of its family tree, i.e. the basic objective of *phylogenetic* analysis. To this end, the comparative method continues to privilege regular sound changes and morphology (especially paradigms, where available). It can seem questionable, then, that many Bayesian phylogenetic approaches are described as using ‘lexical data’ instead. It is easy to misconstrue and overplay, however, the apparent contrast between the comparative method and ‘mere’ lexical data. For a start, the preference for morphological paradigms belongs only to a best-case wish-list. In reality, extensive paradigms are simply not available in all language types. With isolating languages, the comparative method has to make do without them — hence the greater difficulties and doubts surrounding the deep reconstruction of families such as Sino-Tibetan. Lexical data are universally available (even if somewhat more limited in polysynthetic languages). There are also under-documented and extinct languages for which lexical data — short word lists — are all we have, and all we ever will have.

As for sound laws, in order to establish that a sound correspondence is regular and repeated, the comparative method requires multiple tokens of that correspondence. Where are those tokens to be found? Firstly, grammatical morphemes are so few in number that most tokens are necessarily found in the lexicon instead. Secondly, tokens of regular sound changes are by definition found mostly in cognates. (Loanwords may mimic *some* past sound changes when phonologically adapted into the borrower language, but otherwise show only those changes that arose thereafter.) Thirdly, cognate tokens are not evenly distributed across the whole lexicon, but are necessarily concentrated in that part of it least susceptible to lexical replacement (whether by loanwords or language-internal semantic shifts). So while in principle there is no special focus on basic vocabulary when searching for tokens of regular sound correspondences, in practice the comparative method finds much of its evidence (more than in any other comparable sector of the vocabulary) in exactly the same stable core of the lexicon that phylogenetic analyses use.

More importantly still, much of the confusion and dismissal of the phylogenetic value in mere ‘lexical data’ is because that term is in fact a popular misnomer. The crucial comparative linguistic datum used is not lexicon, but cognacy. The 0 and 1 (or multi-state) values that form the input to phylogenetic analyses do not directly represent lexemes at all, or their phonological forms. What the 0s and 1s represent, directly and uniquely, are relationships of cognacy, i.e. common descent. Simply drawing up the Swadesh list for a given language is not enough. The list of lexemes is not the data, and is unusable by the method until each lexeme has been assigned to its correct cognate set *vis-à-vis* the lexemes in all other languages in that family. Only then is there any actual comparative data to be fed into the phylogenetic analysis — because, to repeat, the input data are not lexical forms, but cognacy relationships. So to describe the data as ‘lexical’ is a misnomer particularly because it implies, by omission, no input from phonology (sound correspondences), morphology, and the comparative method in general. Yet these are, on the contrary, all inescapably implicit in the concept of cognacy itself. The task of assigning cognacy, to create the input data, requires detailed, qualitative, historical linguistic analysis, by the comparative method.

Any cognate set is defined by common descent. Specifically, it is defined and identified by the proto-form from which all its members descend directly (i.e. without borrowing, see below). The cognate set to which English *rain* and German *Regen* belong, for example, is defined by the Proto-Germanic form **regna* (www.cobl.info/meaning/rain). Likewise, the fact that Greek *kardia*, Russian *serdce* and Old Irish *cride* are cognate is defined by the Proto-Indo-European form **kerd-*, from which they all descend (www.cobl.info/meaning/heart). That is, to define cognate sets relies on the comparative method, its sound-change laws and reconstructions. Thousands of such individual cognate sets are what provide the phylogenetic signal that Bayesian approaches analyze.

Cognacy assignment, when properly performed, integrates and rests on all of the data, methodology and findings of orthodox comparative-historical linguistics, not least in phonology and morphology. It is only that methodology that can reliably establish which word-forms are true cognates to each other, even when sound changes have obscured original similarities. Likewise, only the comparative method can tell apart true cognates from loanwords or from chance lookalikes, essentially because those do *not* exhibit the expected regular sound correspondences. The data that Bayesian phylogenetic analysis uses are best not described as just ‘lexical data’, then, but more accurately as *cognacy data*, for a sample list of precisely defined *lexical* meanings.

Finally, one should be under no illusions as to the many complexities and challenges in preparing cross-linguistic data-sets of cognacy in lexical meanings. Early data-sets such as that by Dyen, Kruskal & Black (1992) have been rightly criticized for many failings; much progress is still being made towards more rigorous and consistent policies for drawing up such databases, specifically as appropriate for the purposes of phylogenetic analysis and chronology estimation. The methodology needed is too complex to do it justice in this chapter, however, and is set out in much more detail in Heggarty et al. (in prep.). For a major language family, assigning cognacy across over scores or hundreds of language varieties, ancient and modern, in up to 200 target meanings for each, is a daunting task. Even for the best-researched language families like Indo-European, it demands painstaking historical linguistic analysis by large teams of language and family experts. That all serves only to re-emphasize the fundamental point: far from being discarded, all the methodology and findings of orthodox historical linguistics are precisely what databases of *cognacy* in basic lexicon are founded upon in the first place.

4.3 Size Matters

There is one other crucial advantage of data in the form of cognacy in lexical meanings: there are enough of them. For Bayesian phylogenetic analysis to work well requires very significant amounts of data in order to estimate all main aspects of the results: the tree structure, branch lengths, model parameters, the corresponding time-depths, and so on. The large quantities of data needed are only really available from cognacy in the lexicon. The experience of Ringe et al. (2002), in their (non-Bayesian) search for a “perfect phylogeny” for the Indo-European family, is instructive here. Ringe et al. start out with the historical linguist’s instinctive preference for data characters in phonology and morphology. But they find only 22 of the former, and 15 of the latter, that they consider validly usable, and retain even “fewer morphological characters” in their follow-up study (Nakhleh et al. 2005: 394). This not only makes for a very small data-set, but as Ringe et al. (2002: 98) themselves admit, “the worst news is yet to come: the vast majority of our well-behaved monomorphic characters simply define one or more of the ten uncontroversial subgroups of the family, contributing nothing to their higher-order subgrouping”. In practice, the main higher-order nodes in their output trees turn out to be supported by just two, one or even none of these phonological and morphological characters. Most of the phylogenetic information ends up being provided by their lexical cognacy characters after all, thanks to the sheer number of them.

A further consideration is that Ringe et al.’s phonological characters are effectively all just binary. Their P2 character “full ‘satem’ development of dorsals”, for example, allows values of just present or absent. Moreover, any one language cannot change from one of those states to the other more than once over the entire history of Indo-European. Each lexical meaning, however, typically corresponds to a range of many different cognacy states, effectively multiplying the amount of discriminatory data in that one meaning. For a broad family such as Indo-European, the average meaning has well over ten different cognacy states (for exact statistics per meaning, see www.cobl.info/wordlist/Jena175). And in each meaning, changes from one (cognacy) state to another can arise continuously, across the tree, throughout the family’s divergence history. A 200-meaning list thus turns into many thousands of cognacy states, and even if not all of them are informative on the higher-order branching, they all still contribute crucially to the estimation of branch lengths and time-depths.

This is also why phylogenetic analyses are better off aiming for data-sets of the order of Swadesh’s initial 200 meanings, rather than shorter versions slimmed down for other purposes, such as Swadesh’s own 100-meaning list or Tadmor et al.’s (2010: 238-243) Leipzig-Jakarta list, also of just 100 meanings. This need for sheer quantity of data is another reminder of the scale of the task to draw up such comparative databases.

4.4 Not Cherry-Picking, But ‘Safety in Numbers’

Large data-sets, based on a pre-determined list of meanings intended to be applied to any language family, also have the advantage of forcing objectivity. They avoid the danger of individual scholars subjectively cherry-picking data characters, not least in phonology and morphology, that may be unrepresentative and open to interpretation (including, for instance, as to which state was ancestral). It is striking that historical linguistics has still failed to come to a consensus family tree even for many of the best researched of all language families, including Romance, Germanic, and Indo-European as a whole. In part this is because of how the discipline has essentially proceeded on the disputed points: competing scholars each invoke small subsets of the language data, cherry-picked to argue for one tree structure over another, but which fail to convince while other cherry-picked subsets point to rival analyses. Subjectively *excluding* other selected data, meanwhile, comes with a risk of circularity, in that particular characters may be judged unreliable and excluded precisely because they do not fit with a preconceived idea of what the tree structure should be.

Ringe et al.'s (2002) data-set is heavily “screened”, one of the reasons why they end up with so few informative characters, but they are still frustrated by far more characters incompatible with any tree (at least 16) than supporting the main higher-order nodes (mostly 4 or less). It is not that all language data necessarily reflect phylogeny; of course not. But determining which innovations most likely arose in parallel, for instance, and on which different branches, is precisely one of the results that emerges in any case from a powerful Bayesian phylogenetic analysis — and in the form of a balanced, quantitative assessment rather than a subjective and potentially circular judgement. We see shortly below the obvious dangers of cherry-picking data for chronological purposes, too, when trying to second-guess how much or how little language change is ‘plausible’ within a given time-span. Avoiding such subjectivity is certainly one of the attractions of large data-sets of cognacy across a pre-determined list of lexical meanings.

4.5 Why Use the Lexicon for Dating?

Dating is important – it enables us to link linguistic divergences to archaeological and historical events, and build a coherent story about human prehistory. However, it is one thing to establish the phylogeny of a language family; researching its chronology is quite another. And the fact that the discredited technique of glottochronology was directly based on *lexicostatistics* has given the lexicon a bad name for this purpose. The basic problem is not the lexicon itself, however, nor the concept of cognacy. Rather, the undoing of glottochronology was its unrealistically strict assumptions, especially of a universal, unvarying rate of change, and its simplistic distance-based *measures*, without any real *model* of linguistic ‘descent with modification’ through time. Certainly, it is hardly as if historical linguists have enjoyed much success and unanimity in looking to other levels of language, beyond the lexicon, to help in estimating time-depths.

4.6 Dating by Phonetics or Morphosyntax?

In phonology and phonetics, for instance, rates of change and divergence appear to vary even more freely and spectacularly than is usual in the core lexicon. Take for example Latin [akʷam] *water*, [altom] *high* and [ad illom] *to that*, which in only two millennia have all attrited, in French, to just [o] (distinguished only in spelling as <eau>, <haut> and <au>), alongside other spectacular examples of change such as [kalidom]→[fo], [kaballos]→[fo] and [habitom]→[y] (*hot*, *horses* and *had*). Or one can just as well cherry-pick examples of very slow sound change, even over the many millennia since Proto-Indo-European. Reconstructed *bʰréh₂t(ē)r *brother*, for example, still retains its initial obstruent + rhotic cluster little changed into modern Welsh [braut], Russian [brat], English [brʌðə] and ‘even’ French [frɛʁ].

With such great variation in possible rates, arguing just from one-off examples leaves us open to cherry-picking instances where change has been either abnormally fast or abnormally slow, and thus unrepresentative and no proof of anything for dating purposes. Rather, it leaves us open to impressionistic and subjective expressions of conviction — precisely what we need to get away from, in favor of an objective, balanced distribution of rates of change, and an explicit historical model.

A keystone in the traditional short chronology for Indo-European, for instance, is the perception that Avestan and Vedic are ‘so close’ that the divergence between them ‘must’ be but a matter of a few centuries. Sims-Williams (1998: 126), for example, offers selected sentences that “may be transposed from the one language into the other merely by observing the appropriate phonological rules”. Much the same, though, can be said of selected phrases in Italian and Spanish, for instance: see Heggarty & Renfrew (2014: 545). Change and divergence have been minimal in no end of word pairs, such as Italian [lingwa] vs. Spanish [lɛngwa] *tongue* (from Latin [lingʷa]), [mɔndo]-[mundo]

world, [θjelo]~[tʃjelo] *sky*, etc. Indeed, other cases show no real divergence at all: [kanto]~[kanto] / *sing*, [salta]~[salta] *s/he jumps*, and so on. Here too, simply applying phonological rules can straightforwardly transpose one language to the other, but that hardly proves a time-depth of divergence of just a few centuries, as traditionally envisaged between Avestan and Vedic. For the net divergence between Italian and Spanish to arise, out of Latin, took the same two millennia that in French saw much more radical change.

Morphosyntax, likewise, does not seem to offer a viable data-type for language dating. Starting out from the inflectional case system of Proto-Indo-European, for instance, its modern descendants show an entire gamut of different amounts of net case loss, from total to almost nil. At one extreme, French has lost all traces of the system on nouns, while English retains only its ‘Saxon genitive’. Other languages have maintained parts of the system, albeit heavily eroded (e.g. Romanian and German). At the other extreme, meanwhile, most Slavic languages still count up to six cases. So over the same time-depth since Proto-Indo-European, different branches and languages have lost (and sometimes, re-created) cases at very different rates and stages in time. Note how Bulgarian, for example, has bucked the trend of Slavic continuity and seen most of its case system collapse since the time of Proto-Slavic. It seems hard to make anything of this for chronological purposes. Take the famous case of Lithuanian, often evoked in admiration at how little some of its case-endings have changed since Proto-Indo-European. This is of no help to decide between the two main hypotheses on the time-depth of Indo-European: little change over six millennia hardly excludes little change over another few millennia, either. A further consideration is that different levels of language can change at very different rates. Within Romance, French is unquestionably an outlier in how much phonetic attrition it has undergone. In lexis, however, French is unremarkable, a typical Romance language with retention and cognacy rates of exactly the same order as most of its sister languages.

Given all this variability, it is no surprise that those historical linguists who dare to make pronouncements on their ‘intuitions’ of how long a time-span is or is not plausible for a given range of divergence across a family do not necessarily agree with one another. Many Indo-Europeanists, seduced by the claims of linguistic palaeontology (likewise subjective and disputed, see Heggarty 2014, Heggarty forthcoming), have taken as read a Steppe hypothesis time-depth of c. six millennia. That assumption has then set their benchmark for intuitions on what rates of divergence are ‘plausible’, and forced further assumptions on crucial chronological questions such as the time-depth of the split between Indic and Iranian that led to Vedic and Avestan. Remove the assumptions, however, and other linguists such as Dolgopolsky (1990: 239) see the Steppe hypothesis time-frame espoused by Mallory (1989) as “utterly unrealistic”. Mallory’s “dating, which presupposes that Proto-Anatolian, Proto-Indo-Iranian, Greek and other descendant languages could have diverged from each other for a mere 2000 years, is absolutely inconceivable” for Dolgopolsky. Many more linguists would challenge other aspects of Mallory’s (1989) chronology, in which Proto-Germanic, for instance, as recently as 2500 years ago, is “a late Indo-European dialect”.

4.7 Why Linguistic Dating *Should* Be Possible

Whatever levels of language one looks to, then, cherry-picked examples and impressionistic pronouncements from ‘intuition’ lead us to a dead-end, a stand-off between opposing subjective positions. So it is perhaps understandable that some linguists steer clear of any idea of dating from language. “Linguists don’t do dates”, as McMahon & McMahon (2006) put it, while for Dixon (1997: 47, 49) “What has always filled me with wonder is the assurance with which many historical linguists assign a date to their reconstructed proto-language... How do they know? ... It does seem to be a house of cards”.

Such pessimism, however, actually loses sight of some basic tenets of historical linguistics. First, it is not entirely mysterious why rates of change can vary; there are often good explanations, both language-internal and language-external. Sub-systems within a language may long be stable, but then once perturbed (sometimes as a knock-on effect of changes on other levels of the language) can change rapidly until they settle into a new stable system, as in the case-system collapse in Bulgarian, or phonological system change in Old Irish. External contexts, too, can vary from relative isolation to punctuations of intense contact. This is widely seen as the explanation for Bergsland & Vogt's (1962) paradigm case of the contrast between high cognate retention in Icelandic alongside bursts of lexical turnover in English, triggered by the Viking and Norman conquests. In short, where rates depart significantly from 'default' expectations, there is often a good reason.

Second, it is clear that language change is unstoppable. So however much the rates may vary, changes do at least build up progressively through time. Indeed, in practice many of the same historical linguists who repeat the *de rigueur* disclaimers that glottochronology is not to be trusted do themselves go on to invoke its 'dates', for want of anything better; see Kaufman & Golla (2000: 52), for instance, on the major language families of the Americas. So although glottochronology as a method is firmly discredited, on both of the above grounds there remains a widespread, grudging acceptance that rates of change over time are at least not a complete free-for-all. Or to put it more formally: notwithstanding some much-cited instances of exceptional change or stability, those plausibly represent the extremes of a distribution, even if a broad one, of more usual and more modal rates of change.

That already brings us closer to a potential way of making use of these basic tenets of historical linguistics for a much more sophisticated, objective and realistic approach to language dating. Certainly, with all the discipline's knowledge of language history and how it proceeds, we should be able to do far better than just simplistically totting up numbers up to 100 or 200, as glottochronology did. Indeed, instead we can enlist another fundamental tenet, namely the process of linguistic descent with modification by which any family of languages diverges out of its common ancestor. This process is precisely what phylogenetic analyses serve to model explicitly, which is why they are in principle so applicable to language. The mere counts of lexicostatistics and the very simple formula applied to them by glottochronology, meanwhile, have no such explicit model of the actual process of language descent with modification.

4.8 Why Lexical Cognacy for Dating?

Given all those possibilities from Bayesian phylogenetics, it is worth reconsidering which type of language data might be most suitable as input, specifically for chronological purposes. We have already seen some of the problems with other levels of language, and in fact cognate status in the lexicon turns out to have one distinct advantage, specifically for chronological purposes, even over other, form-based data types that can seem more typically what historical linguists would use. For chronology, what is needed is a set of **comparative data points** that are free to change iteratively and repeatedly, so that their 'clock' does not at some point stop ticking and no longer allow change. This can be a problem with most form-based linguistic data: sounds or morphemes that disappear completely generally cannot then reappear, to start changing again. Take the example of ablaut grades in Indo-European. Many Indo-European roots are found shared across multiple branches of Indo-European, but from one branch to the next the lexemes used can be based on different ablaut grades. In the meaning FOOT, for example, Germanic has the lengthened o-grade of PIE *ped- (hence *foot*), Greek *πούς, ποδός* comes from the (short) o-grade, and Latin *pes, pedis* from the e-grade.

Certainly, while a cognate set is essentially defined by a shared root morpheme, it can be helpfully subdivided into its various ‘cognate sub-sets’, according to finer morphological or phonological criteria such as different ablaut grades of that root, or the presence of one or other affix additional to the root. Ultimately, though, the ablaut alternation system of Indo-European largely broke down, and lexemes in different branches fossilized their forms, which were no longer free to change any further in this criterion. By the time of Latin, *pes*, *pedis* was already long fixed as of e-grade, and thereafter all Romance languages inherited that fixed datum. Italian *piède*, or English *foot*, cannot change (back) to any other Indo-European ablaut grade, which ceased being a meaningful system many millennia ago.

From a dating perspective, then, ablaut grade is, unhelpfully, a datum for which the change clock stopped ticking definitively, long ago deep in the history of the Indo-European family. The clock that does remain ticking, meanwhile, is (root) cognate status. Italian and English, like any other Indo-European language, could in principle ultimately switch to a new lexeme for FOOT, in a different (root) cognate set. This ability to iteratively change, for the clock never to stop ticking, is a crucial characteristic that makes cognate status in lexical meanings the data type of choice specifically for the purpose of *dating*. It is also one of the main reasons why it is important and most advantageous to analyze cognate status at the level of the root, rather than to use cognate subsets to try to contribute to the chronological estimation.

Note that there are in fact two tasks, then, that it is possible to take somewhat separately: recovering the phylogeny, and dating that phylogeny. The preference for using only the cognacy turnover data for dating, then, does not limit the language data used to contribute to determining the phylogeny. To that end, one certainly can include finer analysis such as breaking down root cognate sets into their respective subsets, on further morphological and phonological criteria. Indeed, one can make use of morphological and phonological characters that are not iterative, but one-off or sequential changes. These can even be used to constrain the model to return only trees compatible with those characters — but then on that constrained phylogeny one can use only the lexical cognacy data to contribute to the dating. (See Atkinson et al. (2005: 209) for an analysis constrained to Ringe et al.’s (2002) data, including notably their key phonological and morphological criteria.)

4.5 How do Bayesian phylogenetic methods actually date language divergences?

To *calibrate* the clock, historical information is used to constrain the age of known branches. These calibrations allow the analysis to estimate the clock rates in regions where the timing is known, and to extrapolate these to infer rates in regions where the timing is unknown. While these calibrations can be simple point estimates (e.g. a given year), it is more common to specify these as probability distributions, as this allows uncertainty in the divergence time estimates to be explicitly accounted for (Ho and Phillips 2009). For example, Birchall, Dunn, and Greenhill 2016 used this approach to date the origin of the Chapacuran languages in South America to around 1,040 years ago. One of the calibrations they used was based on Meireles 1989’s suggestion that Moré and Cojúbim speakers were a single population that diverged when the Jesuits arrived in the region in the 1740s, as no mention is made of the Cojúbim until 1781. Birchall, Dunn, and Greenhill 2016 implemented this calibration as a log-normal distribution with a 95% probability that Moré and Cojúbim diverged between 213 and 723 years ago — so the lower bound was close to the 1780s date, while the upper bound stretched 500 years in the past to allow for the two communities to have diverged somewhat earlier than the Jesuits’ arrival, in case Meireles 1989 was incorrect.

Given the serious failings of glottochronology, how confident can we be that these methodological advances are accurate? In a simulation test of phylogenetic methods for resolving language history (Greenhill, Currie, and Gray 2009), we evaluated the performance of the rate-smoothing approach. First, we simulated linguistic data on two known phylogenies (described in more detail in section 5.2 below). We then calibrated two arbitrarily chosen nodes on each tree and constrained them to be within $\pm 10\%$ of their true age. We found that the rate smoothing approach was able to estimate a value close to the true root age even when borrowing rates reached up to 15% – the difference between the true age and the estimated age was 2.2% or 6.1% depending on the tree topology. As borrowing increased, the estimated age of the tree decreased. This result shows that these methods are fairly robust to most reasonable levels of borrowing and can reasonably accurately estimate language divergence times. The accuracy and performance of these methods continues to be an ongoing research topic within phylogenetics and performance will likely continue to improve (Drummond and Suchard 2010; Lanfear, Welch, and Bromham 2010; Duchêne, Lanfear, and Ho 2014).

5. How accurate are the results of Bayesian phylolinguistics?

How well do Bayesian phylogenetic methods work at recovering language history? As with any inferential tool, it is critical to assess the performance and accuracy. There are two ways we can assess the performance and accuracy of phylogenetic methods on linguistic data. The first way is to validate the phylolinguistic results with results from the comparative method, while the second validates the results of analyses of simulated data, where the complete history is known.

5.1 Verification with the comparative method.

Many of the studies using phylogenetic methods discuss how consistent their results are with those of the comparative method. Studies showing strong concordance with the comparative method and their phylogenetically estimated subgroups include Ainu (Lee & Hasegawa 2013), Aslian (Dunn et al. 2011), Austronesian (Gray et al. 2009), Bantu (Currie et al. 2013, Grollemund et al. 2015), Chapacuran (Birchall et al. 2015), Huon Peninsula (Greenhill 2015), Indo-European (Bouckaert et al. 2012, Chang et al. 2015), Japonic (Lee & Hasegawa 2014), Pama-Nyungan (Bowerman & Atkinson 2012), Semitic (Kitchen et al. 2009), Timor-Alor-Pantar (Robinson and Holton 2012), Tupian (Galucio et al. 2015), Tupi-Guarani (Michael et al. 2015), to Uralic (Honkola et al. 2013, Syrjänen et al. 2013).

To date perhaps the best comparison between the findings of the traditional comparative method and phylogenetic methods can be found in the Austronesian languages. There are strong hypotheses about the origin and subgrouping of these languages derived from the comparative method, and robust dates for many of the main groups derived from archaeology (Diamond and Bellwood 2003; Bellwood and Dizon 2008; Blust 2013; Pawley 2002; Green 2003; Ko et al. 2014). The main hypothesis holds that the Austronesian languages originated in Taiwan, where 9 of the 10 major branches are still spoken (Blust 1999; Blust 2013) around 5500 years ago (Blust 2013). The remaining branch – Proto-Malayo-Polynesian – spread south through the Philippines around 1000 years later (Bellwood and Dizon 2008; Blust 1999; Pawley 2002), through Indonesia and along the coast of New Guinea. By around 3000-3200 years ago, the Austronesians can be strongly linked to the development of the Lapita cultural complex in Near Oceania (Sheppard et al. 2015; Green 2003) which brought a distinctive dentate stamped and red-slipped pottery style, a range of domesticates, and social practices. Finally, around 2800-3000 years ago the Austronesians settled Western Polynesia, paused again for 1000 years, and then entered East Polynesia (Wilmshurst et al. 2011).

Importantly, Austronesian is a difficult test case as it is a large family with more than 1200 languages that originated from a rapid population expansion into and across a range of new environments, encountering speakers of very different languages, and undergoing substantial shifts in population size. All of these factors have resulted in huge variation in cognate retention rates between languages (Blust 2000; Greenhill 2015). One of the largest lexicostatistical studies ever conducted (Dyen 1962) on these languages produced bizarre results, concluding that the homeland of the Austronesian language family was in Near Oceania rather than in Taiwan (Greenhill et al. 2010).

We used phylogenetic methods on 400 Austronesian languages to test between the above ‘pulse-pause’ settlement scenario and an alternative proposal from genetics. This alternative scenario, the ‘slow boat’, suggests the Austronesians originated in Wallacea (the region around modern day Sulawesi) between 13,000 and 17,000 years ago before spreading north into the Philippines and Taiwan, and east into Oceania (Oppenheimer and Richards 2001; Soares et al. 2008). Our results (Gray et al. 2009) showed overwhelming support for the first ‘pulse-pause’ scenario and none for the ‘slow boat’ scenario. This support for the pulse-pause scenario was both for the broad tree topology and age of the family; we estimated a mean age of 5230 years (95% highest posterior density interval of 4730-5790 years).

In subsequent papers, we carefully evaluated the similarity between the phylogenetic tree and the subgroupings proposed by the comparative method (Greenhill et al. 2010; Greenhill and Gray 2012). Our analysis did in fact recover most of the key subgroupings identified by the comparative method. Figure 4 visualizes this strong consistency between the linguistically attested subgroupings (based on the classification in the *Ethnologue* (Lewis 2009) and the phylogenetic results. The similarities are overwhelming – and statistically significant (Greenhill et al. 2010; Greenhill and Gray 2012). We recovered the following groups (working down the tree): Malayo-Polynesian (Dahl 1973), the Philippine ‘microgroups’ e.g. Bashiic, Central Luzon, Central Philippines, Cordilleran, Manobo, Palawanic, and Subanun (Blust 1991), Malayic (Adelaar 1992), Chamic (Thurgood 1999), Celebic (Mead 2003), Greater South Sulawesi (Adelaar 1994), North Sarawak (Blust 1974), Eastern Malayo-Polynesian (Blust 1978), Central-Eastern Malayo-Polynesian (Blust 1978), Bima-Sumba (Blust 2008), Central Maluku (Collins 1982), Yamdena-North Bomberai (Blust 1993), and South-Halmahera/West New Guinea (Blust 1978). We recover the large Oceanic subfamily [Dempwolff (1927); Ross 1998], and its Near Oceanic component groups including Temotu (Ross and Næss 2007), South-East Solomonic (Pawley 1972), Admiralties (Ross 1988), Papuan Tip (Ross 1988) and Meso-Melanesian (Ross 1988). The majority of the North New Guinea languages are strongly grouped together (Ross 1988) although some languages from Willaumez are misplaced due to unidentified loan words (Greenhill et al. 2010). Remote Oceania is also well resolved with the phylogenies recovering Central Pacific (Grace 1959), Polynesian [Pawley (1966); Marck 2000], Micronesian (Bender et al. 2003), South Vanuatu (Lynch 2001), North and Central Vanuatu [Tryon (1976); Clark 2009], New Caledonia and the Loyalties (Ozanne-Rivierre 1992).

To be fair, there are some mismatches with the expected groupings. In total, 25 of the 400 languages are misplaced (Greenhill, Drummond, and Gray 2010). Some of these misplacements are due to errors in the data – e.g. the Willaumez languages (Nakanai, Maututu, Lakalai) sit closer to the North New Guinea languages than to their expected sisters in the Meso-Melanesian subgroup. This misplacement is probably due to unidentified lexical borrowings between the Willaumez languages

and their Meso-Melanesian neighbours in west New Britain¹. However, many of the other misplacements reflect long-standing classification difficulties. For example, it is unclear whether Irarutu should be in the Central Malayo-Polynesian or the South-Halmahera/West New Guinea subgroup (Blust 1993). Another group we do not recover is Central Malayo-Polynesian, but this group is now thought to be a dialect chain with low internal cohesion, and is only supported by overlapping isoglosses (Blust 1993). So while there are some errors in the data – discussed in full in Greenhill et al. (2010) – the Bayesian phylogenetic trees do remarkably well at recovering both the known subgroups and the places of uncertainty.

5.2 Validation with simulations

The second way we might validate phylogenetic inference is to simulate data on a known ‘true’ phylogeny and then analyze the simulated data to see if we recover the ‘true’ tree. This approach is commonly used in the biological phylogenetics literature to test the effectiveness of these methods. Greenhill et al. (2009) applied this logic to test how well Bayesian phylogenetic methods could recover the true tree with lexical cognate data. First, we created two tree topologies: one shaped like Polynesian with an unbalanced, chained topology, and the other shaped like Uto-Aztecan with a well-balanced tree. Then we used a simple ‘stochastic Dollo’ model of linguistic evolution (Nicholls and Gray 2006; Atkinson et al. 2005) to simulate cognate data. This model started at the root of the tree and worked its way down to the tips evolving new cognate sets to create a new simulated dataset.

One of the major objections to phylogenetic methods in linguistics is that the borrowing of items between languages invalidates the tree (Terrell 1988, Moore 1994). Rather than throwing methods away at the first sign of trouble it is better to assess performance and identify at what point the method breaks down. To this end we also simulated borrowing events in these data. During the simulation process, traits were randomly selected and copied to another lineage to simulate the borrowing of a cognate set between languages. We varied the amount of borrowing from 0% to 50% of the total cognates.

Finally, for each of the simulated datasets, we used Bayesian phylogenetic methods to estimate the posterior probability distribution. For each analysis we then constructed the maximum clade credibility tree (a single summary tree of the posterior) and then compared how close this estimate was to the ‘true’ tree. Our results show that when there is no borrowing the analysis finds a tree that statistically is almost indistinguishable from the true tree. As the amount of borrowing in the data increased, so did the difference between the true tree and the tree recovered by the phylogenetic analysis. What level of borrowing is plausible? A fairly extreme example of language borrowing is English (leaving aside creoles and mixed languages), which has borrowed more than 60% of its total lexicon from French and Latin (Embleton 1986). However, only 16% of English’s basic vocabulary is borrowed (Embleton 1986). Given that most language phylogenies are built from data on cognacy in basic vocabulary, we took 0-20% as the extreme range of *undetected borrowing* in the type of datasets used by language phylogenies. Over this range, the difference between the true trees and the recovered trees was very small – the balanced topology showed an average perturbation of at most 0.7%, while the more fragile chained topology showed an average perturbation of 6.8% at most. In summary, this simulation study shows that Bayesian phylogenetic

¹ In fact, Ross no longer claims that the Willaumez languages form a valid subgroup, see “Western Oceanic revisited” (unpublished draft, cited with permission).

methods are exceptionally good at finding a tree very close to the true history, even in the face of realistic levels of undetected borrowing.

6. Conclusion

The view we have taken here is that Bayesian phylogenetic methods are a powerful supplement to traditional linguistic scholarship — not a replacement. These approaches look set to become part of orthodox methodology in historical linguistics, as a powerful new tool to complement the tried-and-tested qualitative analyses. In combination with the comparative method, they enable us to estimate the strength of support for competing subgrouping proposals, as already applied to a dozen or so language families from across the world (see section 5.1 for references). Further methodological developments promise to enrich and expand these analyses. In the near future we expect to see more combined analyses of lexical, phonological and morphological data. We also look forward to the development of increasingly sophisticated models of cognate evolution and tools for "computer assisted" cognate detection.

Bayesian phylogenetic methods can also be used for far more than just inferring language trees and dating them. They lend themselves to being combined with geographical models too (phylogeography), and have thus contributed to the debates on the homelands of Bantu, Indo-European, Japonic and Uralic (Grollemund et al. 2015, Bouckaert et al. 2012, Lee and Hasegawa 2014, Honkola et al. 2013). The ability of Bayesian phylogenetic tools to robustly quantify branch lengths and infer divergence dates can also be applied to identify the relative stability of both lexical and grammatical items (Pagel et al. 2007, Dediu 2011) and to investigate the forces shaping these rates of change (Pagel et al. 2007, Greenhill et al. 2017). Other work has explored how grammatical systems have evolved — including aspects such as word-order (Dunn et al. 2011), argument structures (Dunn et al. 2017) and numeral systems (Calude & Verkerk 2016) — and has shown that languages have a persistent evolutionary bias against ergativity (Bickel et al. 2015). Finally, these tools have been used to tackle questions of linguistic anthropology, from the evolution of color term systems in Australia (Haynie and Bower 2016) to kinship system terminology (e.g. Jordan 2011, Guillon and Mace 2016) and the rise and fall of political complexity in the Pacific (Currie et al. 2010).

As this diverse and growing range of illustrations shows, Bayesian phylogenetics offers linguists an enticing opportunity to make so much more of their hard-won language data and their two centuries of qualitative research into language histories. Language phylogenies can be recruited as reference frameworks of 'descent with modification' in other characteristics of language and culture. Bayesian modelling, in particular, holds out the prospect that historical linguistics can at last realize to the full its potential to contribute not just to other branches within linguistics, such as syntax and universals, but also beyond, to other disciplines such as anthropology and cultural history. If historical linguists can exorcise the ghost of lexicostatistics past, then there is an exciting future to embrace.

References

- Adelaar, K. Alexander 1992. *Proto-Malayic: A Reconstruction of Its Phonology and Part of Its Morphology and Lexicon*. Canberra: Pacific Linguistics.
- Adelaar, K. Alexander 1994. "The Classification of the Tamanic Languages." In *Language Contact and Change in the Austronesian World*, edited by T. Dutton and D. Tryon, 1–42. Berlin: Mouton de Gruyter.
- Atkinson, Quentin D. and Gray, Russell D. 2005. "Curious parallels and curious connections: Phylogenetic thinking in biology and historical linguistics." *Systematic Biology*, 54(4), 513–526.
- Atkinson, Quentin D., Geoff Nicholls, David Welch & Russell D. Gray. 2005. "From words to dates: water into wine, mathemagic or phylogenetic inference?" *Transactions of the Philological Society* 103(2): 193–219. <http://dx.doi.org/10.1111/j.1467-968X.2005.00151.x>
- Bellwood, Peter, and Eusebio Dizon. 2008. "Austronesian Cultural Origins: Out of Taiwan, via the Batanes Islands, and Onwards to Western Polynesia." In *Past Human Migrations in East Asia: Matching Archaeology, Linguistics, and Genetics*, edited by A Sanchez-Mazas, R. Blench, M. D. Ross, I. Peiros, and M. Lin, 23–39. London: Routledge.
- Bender, Byron. W., Ward H. Goodenough, F. H. Jackson, Jeff C. Marck, Kenneth L. Rehg, H. Sohn, S. Trussel, and J. W. Wang. 2003. "Proto-Micronesian Reconstruction." *Oceanic Linguistics* 1 (42): 272–358.
- Bergsland, Knut, & Hans Vogt. 1962. "On the validity of glottochronology." *Current Anthropology* 3(2): p.115–153. www.jstor.org/stable/2739527
- Bickel, Balthasar, Johanna Nichols, Taras Zakharko, Alena Witzlack-Makarevich, Kristine Hildebrandt, Michael Rießler, Lennart Bierkandt, Fernando Zúñiga & John B. Lowe. (2017). *The AUTOTYP typological databases*. Version 0.1.0 <https://github.com/autotyp/autotyp-data/tree/0.1.0>
- Bickel Balthasar, Alena Witzlack-Makarevich, Kamal K. Choudhary, Matthias Schlesewsky and Ina Bornkessel-Schlesewsky I (2015) "The Neurophysiology of Language Processing Shapes the Evolution of Grammar: Evidence from Case Marking." *PLoS ONE* 10(8): e0132819. <https://doi.org/10.1371/journal.pone.0132819>
- Birchall, Joshua, Michael Dunn, and Simon J. Greenhill. 2016. "A Combined Comparative and Phylogenetic Analysis of the Chapacuran Language Family." *International Journal of American Linguistics* 82 (3): 255–84.
- Blust, Robert A. 1974. *The Proto-North Sarawak Vowel Deletion Hypothesis*. University of Hawaii.
- Blust, Robert A. 1978. "Eastern Malayo-Polynesian: A Subgrouping Argument." In *Second International Conference on Austronesian Linguistics: Proceedings, Fascicle I, Western Austronesian*, edited by S.A. Wurm and L. Carrington, 181–234. Canberra: Pacific Linguistics.
- Blust, Robert A. 1991. "The Greater Central Philippines Hypothesis." *Oceanic Linguistics* 30 (73): 129.
- Blust, Robert A. 1993. "Central and Central-Eastern Malayo-Polynesian." *Oceanic Linguistics* 32: 241–93.
- Blust, Robert A. 1999. "Subgrouping, Circularity and Extinction: Some Issues in Austronesian Comparative Linguistics." In *Selected Papers from the Eighth International Conference on Austronesian Linguistics*, edited by Zeitoun E. and P J-K. Li, 31–94. Taipei, Taiwan: Symposium Series of the Institute of Linguistics, Academia Sinica.

- Blust, Robert A. 2000. "Why Lexicostatistics Doesn't Work: The 'Universal Constant' Hypothesis and the Austronesian Languages." In *Time Depth in Historical Linguistics*, edited by C. Renfrew, A. McMahon, and L. Trask, 311–31. Cambridge: McDonald Institute for Archaeological Research.
- Blust, Robert. 2008. "Is There a Bima-Sumba Subgroup?" *Oceanic Linguistics* 47 (1): 45–113. doi:[10.1353/ol.0.0006](https://doi.org/10.1353/ol.0.0006).
- Blust, Robert. 2013. *The Austronesian Languages*. Revised Ed. Canberra: Asia-Pacific Linguistics.
- Bouckaert, Remco R. 2010. "DensiTree: making sense of sets of phylogenetic trees." *Bioinformatics*. 26(10):1372-1373. doi: 10. 1093/bioinformatics/btq110. Epub 2010 Mar 12.
- Bouckaert, Remco R., Philippe Lemey, Michael Dunn, Simon J. Greenhill, Alexander V. Alekseyenko, Alexei J. Drummond, Russell D. Gray, Marc A. Suchard, and Quentin D. Atkinson. 2012. "Mapping the Origins and Expansion of the Indo-European Language Family." *Science* 337 (6097): 957–60. doi:[10.1126/science.1219669](https://doi.org/10.1126/science.1219669).
- Bowern, Claire, and Quentin D. Atkinson. 2012. "Computational Phylogenetics and the Internal Structure of Pama-Nyungan." *Language* 88 (4): 817–45.
- Bowern C. 2016. "Chirila: Contemporary and Historical Resources for the Indigenous Languages of Australia." *Lang Doc Conserv* 10:1–44.
- Burnham, Kenneth P. & David R. Anderson. 1998. *Model Selection and Inference: A practical information-theoretic approach*. New York: Springer.
- Calude, Andrea S. and Annemarie Verkerk (2016) "The typology and diachrony of higher numerals in Indo-European: a phylogenetic comparative study." *J Lang Evol*, 1: 91-108.
- Chang, Will, Chundra Cathcart, David Hall, and Andrew Garrett. 2015. "Ancestry-Constrained Phylogenetic Analysis Supports the Indo-European Steppe Hypothesis." *Language* 91 (1): 194–244. doi:[10.1353/lan.2015.0005](https://doi.org/10.1353/lan.2015.0005).
- Chrétien, C. Douglas. 1943. "The quantitative method for determining linguistic relationships." *University of California Publications in Linguistics*, 1(2): 11-20.
- Collins, James T. 1982. "Linguistic Research in Maluku: A Report of Recent Fieldwork." *Oceanic Linguistics* 21: 73–146.
- Currie, Thomas E, Andrew Meade, Myrtille Guillon and Ruth Mace(2013) "Cultural phylogeography of the Bantu Languages of sub-Saharan Africa." *Proc R Soc B Biol Sci* 280(1762):20130695.
- Dahl, Otto C. 1973. *Proto-Austronesian*. Vol. 15. Sweden: Scandinavian Institute of Asian Studies Monograph Series.
- Dediu, Dan (2011) "A Bayesian phylogenetic approach to estimating the stability of linguistic features and the genetic biasing of tone." *Proc R Soc B Biol Sci* 278(1704):474–479.
- Dempwolff, Otto. 1927. "Das Austronesische Sprachgut in Den Melanesischen Sprachen." *Folia Ethnogllossica* 3: 32–43.
- Diamond, Jared M. and Peter Bellwood. 2003. "Farmers and Their Languages: The First Expansions." *Science* 300 (5619): 597–603. doi:[10.1126/science.1078208](https://doi.org/10.1126/science.1078208).
- Dixon, Robert M.W. 1997. *The Rise and Fall of Languages*. Cambridge: Cambridge University Press.
- Dolgopolsky, Aaron B. 1990. "More about the Indo-European homeland problem." *Mediterranean Language Review* 6–7: p.230–248.

- Drummond, Alexei J., and Marc A. Suchard. 2010. "Bayesian Random Local Clocks, or One Rate to Rule Them All." *BMC Biology* 8 (114): 114. doi:[10.1186/1741-7007-8-114](https://doi.org/10.1186/1741-7007-8-114).
- Drummond, Alexei J., Simon Y. W. Ho, Matthew J. Phillips, and Andrew Rambaut. 2006. "Relaxed Phylogenetics and Dating with Confidence." *PLOS Biology* 4 (5): e88.
- Duchêne, Sebastián, Robert Lanfear, and Simon Y W Ho. 2014. "The Impact of Calibration and Clock-Model Choice on Molecular Estimates of Divergence Times." *Molecular Phylogenetics and Evolution* 78 (June): 277–89. doi:[10.1016/j.ympev.2014.05.032](https://doi.org/10.1016/j.ympev.2014.05.032).
- Dunn, Michael, Niclas Burenhult, Nicole Kruspe, Sylvia Tufvesson, and Neele Becker. 2011. "Aslian Linguistic Prehistory: A Case Study in Computational Phylogenetics." *Diachronica* 28 (3): 291–323. doi:[10.1075/dia.28.3.01dun](https://doi.org/10.1075/dia.28.3.01dun).
- Dunn Michael, Simon J. Greenhill, Stephen C. Levinson and Russell D. Gray (2011) "Evolved structure of language shows lineage-specific trends in word-order universals." *Nature* 473(7345):79–82.
- Dunn Michael, Tonya Kim Dewey, Carlee Arnett, Thórhallur Eythórsson and Jóhanna Barðdal (2017) "Dative sickness : A phylogenetic analysis of argument structure evolution in Germanic." *Language* 93(1):e1–e22.
- Dyen, Isidore. 1962. "The Lexicostatistical Classification of the Malayopolynesian Languages." *Language* 38: 38–46.
- Dyen, Isidore, Kruskal, Joseph B. and Black, Paul. 1992. *An Indoeuropean Classification: A Lexicostatistical Experiment*. Philadelphia: American Philosophical Society.
www.wordgumbo.com/ie/cmp/iedata.txt
- Embleton, Sheila. 1986. *Statistics in Historical Linguistics*. Bochum: Studienverlag Brockmeyer.
- Felsenstein, Joseph. 1978. "The number of evolutionary trees." *Systematic Zoology* 27: 27–33
- Galucio, Ana Vilacy, Sérgio Meira, Joshua Birchall, Denny Moore, Nilson Gabas Júnior, Sebastian Drude, Luciana Storto, Gessiane Picanço, and Carmen Reis Rodrigues. 2015. "Genealogical Relations and Lexical Distances Within the Tupian Linguistic Family." *Boletim Do Museu Paraense Emilio Goeldi: Ciencias Humanas* 10 (2): 229–74. doi:[10.1590/1981-81222015000200004](https://doi.org/10.1590/1981-81222015000200004).
- Grace, George W. 1959. *The Position of the Polynesian Languages Within the Austronesian (Malayo-Polynesian) Language Family. Memoir 16 of the International Journal of American Linguistics*. Indiana: Indiana University Publications in Anthropology & Linguistics.
- Gray, Russell D., Alexei J. Drummond, and Simon J. Greenhill. 2009. "Language Phylogenies Reveal Expansion Pulses and Pauses in Pacific Settlement." *Science* 323 (5913): 479–83. doi:[10.1126/science.1166858](https://doi.org/10.1126/science.1166858).
- Gray, Russell D., Simon J. Greenhill, and David Bryant. 2010. "On the shape and fabric of human history." *Philosophical Transactions of the Royal Society London, B*, 365, 3923–3933.
- Gray, Russell D. & Joseph Watts. 2017. "Cultural macroevolution matters." *Proceedings of the National Academy of Sciences*, 114 (30) 7846–7852.
- Green, Roger C. 2003. "The Lapita Horizon and Traditions-Signature for One Set of Oceanic Migrations." In *Pacific Archaeology: Assessments and Anniversary of the First Lapita Excavation (July 1952)*, edited by C. Sand, 95–120. New Caledonia.
- Greenhill, Simon J. 2015. "TransNewGuinea.org: An Online Database of New Guinea Languages." *PLoS ONE* 10 (10): 1–17. doi:[10.1371/journal.pone.0141563](https://doi.org/10.1371/journal.pone.0141563).

- Greenhill, Simon J., and Russell D. Gray. 2009. "Austronesian language phylogenies: myths and misconceptions about Bayesian computational methods." In Alexander K. Adelaar & Andrew Pawley (Eds) *Austronesian historical linguistics and culture history: a festschrift for Robert Blust*. Canberra: Pacific Linguistics.
- Greenhill, Simon J. and Russell D. Gray. 2012. "Basic Vocabulary and Bayesian Phylolinguistics: Issues of Understanding and Representation." *Diachronica* 29 (4): 523–37. doi:[10.1075/dia.29.4.05gre](https://doi.org/10.1075/dia.29.4.05gre).
- Greenhill, Simon J. Thomas E. Currie, and Russell D. Gray. 2009. "Does Horizontal Transmission Invalidate Cultural Phylogenies?" *Proceedings of the Royal Society, B. Biological Sciences* 276 (1665): 2299–2306. doi:[10.1098/rspb.2008.1944](https://doi.org/10.1098/rspb.2008.1944).
- Greenhill, Simon J., Alexei J Drummond, and Russell D. Gray. 2010. "How Accurate and Robust Are the Phylogenetic Estimates of Austronesian Language Relationships?" *PLoS ONE* 5 (3): e9573. doi:[10.1371/journal.pone.0009573](https://doi.org/10.1371/journal.pone.0009573).
- Greenhill, Simon J., Robert A. Blust, Russell D. Gray. 2008. "The Austronesian Basic Vocabulary Database: From bioinformatics to lexomics." *Evol Bioinform Online* 4:271–283.
- Greenhill Simon J., Chieh-His Wu, Xia Hua, Michael Dunn, Stephen C. Levinson & Russell D. Gray (2017) Evolutionary dynamics of language systems. *Proc Natl Acad Sci*, 114: E8822-E8829.
- Grollemund, Rebecca. 2012. "Nouvelles Approches en Classification: Application aux Langues Bantu du Nord-Ouest." PhD thesis, CNRS/Université Lyon 2, France.
- Grollemund, Rebecca, Simon Branford, Koen Bostoen, Andrew Meade, Chris Venditti, and Mark Pagel. 2015. "Bantu Expansion Shows Habitat Alters the Route and Pace of Human Dispersals." *Proceedings of the National Academy of Sciences of the USA*. doi:[10.1073/pnas.1503793112](https://doi.org/10.1073/pnas.1503793112).
- Guillon Myrtille and Ruth Mace (2016) "A phylogenetic comparative study of Bantu kinship terminology finds limited support for its co-evolution with social organisation." *PLoS One* 11(3):1–15.1.
- Haspelmath, Martin. 2005. *The World Atlas of Language Structures*. Oxford: Oxford University Press.
- Haynie, Hannah J and Claire Bower (2016) "Phylogenetic approach to the evolution of color term systems." *Proc Natl Acad Sci* 113(48):13666–13671.
- Heggarty, Paul. 2010. "Beyond lexicostatistics: How to get more out of 'word list' comparisons." *Diachronica* 27(2): p.301–324. <http://doi.org/10.1075/dia.27.2.07heg>
- Heggarty, Paul. 2014. "Prehistory through language and archaeology." In Claire Bower & Bethwyn Evans (eds) *Routledge Handbook of Historical Linguistics*, 598–626. London: Routledge. <https://www.academia.edu/3687718>
- Heggarty, Paul. forthcoming. "Why Indo-European? Clarifying cross-disciplinary misconceptions on farming vs. pastoralism." *Journal of Indo-European Studies — special issue on Indo-European and Farming*, edited by Gus Kroonen & Bernard Comrie.
- Heggarty, Paul, & Renfrew, Colin. 2014. "South and Island South-East Asia: Languages." In C. Renfrew & P. Bahn (eds) *The Cambridge World Prehistory*, 534–558. Cambridge: Cambridge University Press.
- Ho, Simon Y. W., and Matthew J. Phillips. 2009. "Accounting for Calibration Uncertainty in Phylogenetic Estimation of Evolutionary Divergence Times." *Systematic Biology* 58 (3): 367–80. doi:[10.1093/sysbio/syp035](https://doi.org/10.1093/sysbio/syp035).

- Honkola, Terhi, Outi Vesakoski, Kalle Korhonen, Jyri Lehtinen, Kaj Syrjänen and Niklas Wahlberg (2013) "Cultural and climatic changes shape the evolutionary history of the Uralic languages." *J Evol Biol*:1–10.
- Hymes, Dell. 1960. "Lexicostatistics so Far." *Current Anthropology* 1 (1): 3–44.
- Jordan Fiona M (2011) "A Phylogenetic Analysis of the Evolution of Austronesian Sibling Terminologies A Phylogenetic Analysis of the Evolution of Austronesian." *Hum Biol* 83(2):297–321.
- Kaufman, Terrence., & Victor Golla. 2000. "Language groupings in the New World: their reliability and usability in cross-disciplinary studies." In C. Renfrew (ed.) *America Past, America Present: Genes and Languages in the Americas and Beyond*, 47–57. Cambridge: McDonald Institute for Archaeological Research.
- Kimura, Motoo. 1968. "Evolutionary Rate at the Molecular Level." *Nature* 217: 624–26.
- Kirsch, John A. W. 1969. "Serological Data and Phylogenetic Inference: The Problem of Rates of Change." *Systematic Zoology* 18 (3): 296–311.
- Kitchen, Andrew, Christopher Ehret, Shiferaw Assefa, and Connie J Mulligan. 2009. "Bayesian Phylogenetic Analysis of Semitic Languages Identifies an Early Bronze Age Origin of Semitic in the Near East." *Proceedings of the Royal Society B: Biological Sciences* 270 (1668): 2703–10. doi:[10.1098/rspb.2009.0408](https://doi.org/10.1098/rspb.2009.0408).
- Ko, Albert Min Shan, Chung Yu Chen, Qiaomei Fu, Frederick Delfin, Mingkun Li, Hung Lin Chiu, Mark Stoneking, and Ying Chin Ko. 2014. "Early Austronesians: Into and Out of Taiwan." *American Journal of Human Genetics* 94 (3). The American Society of Human Genetics: 426–36. doi:[10.1016/j.ajhg.2014.02.003](https://doi.org/10.1016/j.ajhg.2014.02.003).
- Kroeber, Alfred L. 1955. "Linguistic Time Depth Results so Far and Their Meaning." *International Journal of American Linguistics* 21 (2): 91–104.
- Lanfear, Robert, John J Welch, and Lindell Bromham. 2010. "Watching the Clock: Studying Variation in Rates of Molecular Evolution Between Species." *Trends in Ecology & Evolution* 25 (9). Elsevier Ltd: 495–503. doi:[10.1016/j.tree.2010.06.007](https://doi.org/10.1016/j.tree.2010.06.007).
- Lee, Sean, and Toshikazu Hasegawa (2013) "Evolution of the Ainu language in space and time." *PLoS One* 8(4):e62243.
- Lee, Sean, and Toshikazu Hasegawa (2014) "Oceanic barriers promote language diversification in the Japanese Islands." *J Evol Biol* 27(9):1905–12.
- Lees, Robert B. 1953. "The Basis of Glottochronology." *Language* 29 (2): 113–27.
- Lewis, Paul M., ed. 2009. *Ethnologue: Languages of the World*. 16th ed. Dallas, Texas: SIL International.
- List, Johann-Mattis, Jananan Sylvestre Pathmanathan, Philippe Lopez and Eric Baptiste. 2016. "Unity and disunity in evolutionary sciences: process-based analogies open common research avenues for biology and linguistics." *Biology Direct* 11(39). 1–17.
- List, Johann-Mattis, Simon J. Greenhill, and Russell D. Gray. 2017. "The Potential of Automatic Word Comparison for Historical Linguistics." *PLoS One*, e0170046. doi:[10.1371/journal.pone.0170046](https://doi.org/10.1371/journal.pone.0170046).
- Lynch, John 2001. *The Linguistic History of Southern Vanuatu*. Canberra: Pacific Linguistics.
- Mallory, J.P. 1989. *In Search of the Indo-Europeans*. London: Thames & Hudson.

- Maurits, Luke and Thomas L. Griffiths (2014) "Tracing the roots of syntax with Bayesian phylogenetics." *Proc Natl Acad Sci U S A* 111(37):13576–81.
- McMahon, April M.S. and McMahon, Robert. 2006. "Why linguists don't do dates." In P. Forster & C. Renfrew (eds) *Phylogenetic Methods and the Prehistory of Languages*, 153–160. Cambridge: McDonald Institute for Archaeological Research.
- Mead, David. 2003. "Evidence for a Celebic Supergroup." In *Issues in Austronesian Historical Phonology*, edited by John Lynch, 115–41. Canberra: Pacific Linguistics.
- Meireles, Denise Maldi. 1989. *Guardiões da Fronteira: Rio Guaporé, século XVIII*. Petrópolis: Vozes.
- Michael, Lev, Natalia Chousou-Polydouri, Keith Bartolomei, Erin Donnelly, Vivian Wauters, and Zachary O Hagan. 2015. "A Bayesian Phylogenetic Classification of Tupi-Guarani." *LIAMES* 15 (2): 1–36.
- Moore, John H. 1994. "Putting Anthropology Back Together Again: The Ethnogenetic Critique of Cladistic Theory." *American Anthropologist* 96 (4): 925–48.
- Moran Steven, Daniel McCloy, & Richard Wright. (2014) *PHOIBLE Online* (Max Planck Institute for Evolutionary Anthropology, Leipzig).
- Murdock, George Peter. 1964. "Genetic Classification of the Austronesian Languages: A Key to Oceanic Culture History." *Ethnology* 3 (2): 117–26. doi:[10.2307/3772706](https://doi.org/10.2307/3772706).
- Nakhleh, Luay, Don Ringe, & Tandy Warnow. 2005. "Perfect phylogenetic networks: a new methodology for reconstructing the evolutionary history of natural languages." *Language* 81(2): 382–420. www.jstor.org/stable/4489897
- Nascimento, F.F., M.D. Reis and Ziheng Yang. (2017). "A biologist's guide to Bayesian phylogenetic analysis." *Nature Ecology & Evolution*, 1, 1446–1454
- Nicholls, Geoff K. and Russell D. Gray. 2006. "Dated Ancestral Trees from Binary Trait Data." *Journal of the Royal Statistical Society: Series B (Statistical Methodology)* 70 (3): 545–66. doi:[10.1111/j.1467-9868.2007.00648.x](https://doi.org/10.1111/j.1467-9868.2007.00648.x).
- Oppenheimer, Stephen J. and Martin B Richards. 2001. "Fast Trains, Slow Boats, and the Ancestry of the Polynesian Islanders." *Science Progress* 84 (3): 157–81.
- Ozanne-Rivierre, François. 1992. "The Proto-Oceanic Consonantal System and the Languages of New Caledonia." *Oceanic Linguistics* 31: 191–207.
- Pagel Mark, Quentin D. Atkinson and Andrew Meade A (2007) "Frequency of word-use predicts rates of lexical evolution throughout Indo-European history." *Nature* 449(7163):717–720.
- Pawley, Andrew. 1966. "Polynesian Languages: A Subgrouping Based on Shared Innovations in Morphology." *Journal of the Polynesian Society* 75 (1): 39–64.
- Pawley, Andrew. 1972. "On the Internal Relationships of Eastern Oceanic Languages." In *Studies in Oceanic Culture History*, edited by Roger C. Green and M. Kelly, 13:3–106. Honolulu: Bernice P. Bishop Museum.
- Pawley, Andrew. 2002. "The Austronesian Dispersal: Languages, Technologies and People." In *Examining the Farming/Language Dispersal Hypothesis*, edited by Peter Bellwood and Colin Renfrew, 251–73. Cambridge: McDonald Institute for Archaeological Research.
- Penny, David, Bennett J McComish, Michael A Charleston, and Michael D Hendy. 2001. "Mathematical Elegance with Biochemical Realism: The Covarion Model of Molecular Evolution." *Journal of Molecular Evolution* 53 (6): 711–23.

- Ringe, Don A, Tandy Warnow, & Anne Taylor. 2002. "Indo-European and computational cladistics." *Transactions of the Philological Society* 100(1): p.59–129. <http://dx.doi.org/10.1111/1467-968X.00091>
- Robinson, Laura C. and Gary Holton. 2012. "Internal Classification of the Alor-Pantar Language Family Using Computational Methods Applied to the Lexicon." *Language Dynamics and Change* 2: 1–27. doi:[10.1163/22105832-20120201](https://doi.org/10.1163/22105832-20120201).
- Ross, Malcolm. 1988. *Proto Oceanic and the Austronesian Languages of Western Melanesia*. Canberra: Australian National University; Pacific Linguistics.
- Ross, Malcolm, and Åshild Næss. 2007. "An Oceanic Origin for Äiwoo, a Language of the Reef Islands." *Oceanic Linguistics* 46: 456–98.
- Sanderson, Michael J. 1997. "A Nonparametric Approach to Estimating Divergence Times in the Absence of Rate Constancy." *Molecular Biology and Evolution* 14 (12): 1218–31.
- Sheppard, Peter J. Scarlett Chiu, and Richard Walter. 2015. "Re-Dating Lapita Movement into Remote Oceania." *Journal of Pacific Archaeology* 6 (1): 26–36.
- Sicoli, Mark A. and Gary Holton. 2014. "Linguistic Phylogenies Support Back-Migration from Beringia to Asia." *PloS One* 9 (3): e91722. doi:[10.1371/journal.pone.0091722](https://doi.org/10.1371/journal.pone.0091722).
- Sims-Williams, Patrick. 1998. Genetics, linguistics, and prehistory: thinking big and thinking straight. *Antiquity* 72(277): p.505–527.
- Soares, Pedro, Jean Alain Trejaut, Jun-Hun Loo, Catherine Hill, Maru Mormina, Chien-Liang Lee, Yao-Ming Chen, Georgi Hudjashov, Peter Forster, Vincent Macaulay, David Bulbeck, Stephen Oppenheimer, Marie Lin & Martin B. Richards. 2008. "Climate Change and Postglacial Human Dispersals in Southeast Asia." *Molecular Biology and Evolution* 25 (6): 1209–18.
- Sokal, Robert R. and Charles D. Michener. 1958. "A Statistical Method for Evaluating Systematic Relationships." *The University of Kansas Science Bulletin* 38 (22): 1409–38.
- Swadesh, Morris. 1950. "Salish Internal Relationships." *International Journal of American Linguistics* 16 (4): 157–67. doi:[10.1086/464084](https://doi.org/10.1086/464084).
- Swadesh, Morris. 1952. "Lexico-Statistic Dating of Prehistoric Ethnic Contacts." *Proceedings of the American Philosophical Society* 96 (4): 452–63.
- Swadesh, Morris. 1955. "Towards Greater Accuracy in Lexicostatistic Dating." *International Journal of American Linguistics* 21 (2): 121–37.
- Syrjänen, Kaj, Terhi Honkola, Kalle Korhonen, Jyri Lehtinen, Outi Vesakoski, and Niklas Wahlberg. 2013. "Shedding More Light on Language Classification Using Basic Vocabularies and Phylogenetic Methods: A Case Study of Uralic." *Diachronica* 30 (3): 323–52. doi:[10.1075/dia.30.3.02syr](https://doi.org/10.1075/dia.30.3.02syr).
- Tadmor, Uri, Martin Haspelmath, & Bradley Taylor. 2010. "Borrowability and the notion of basic vocabulary." *Diachronica* 27(2): p.226–246. <http://dx.doi.org/10.1075/dia.27.2.04tad>
- Terrell, John. 1988. "History as a Family Tree, History as an Entangled Bank: Constructing Images and Interpretations of Prehistory in the South Pacific." *Antiquity* 62: 642–57.
- Thurgood, Graham 1999. *From Ancient Cham to Modern Dialects*. Hawaii: Oceanic Linguistics Special Publications (28).
- Tryon, Darryl T. 1976. *New Hebrides Languages: An Internal Classification*. Canberra: Pacific Linguistics.

- Welch, John J, and Lindell Bromham. 2005. "Molecular Dating When Rates Vary." *Trends in Ecology & Evolution* 20 (6): 320–7. doi:[10.1016/j.tree.2005.02.007](https://doi.org/10.1016/j.tree.2005.02.007).
- Wilmshurst, Janet M, Terry L Hunt, Carl P Lipo, and Atholl J Anderson. 2011. "High-Precision Radiocarbon Dating Shows Recent and Rapid Initial Human Colonization of East Polynesia." *Proceedings of the National Academy of Sciences of the United States of America* 108 (5): 1815–20. doi:[10.1073/pnas.1015876108](https://doi.org/10.1073/pnas.1015876108).
- Yang, Ziheng 1993. "Maximum-Likelihood Estimation of Phylogeny from Dna Sequences When Substitution Rates Differ over Sites." *Molecular Biology and Evolution* 10 (6): 1396–1401.
- Zuckerkandl, Émile, & Linus Pauling. 1965. "History of Evolutionary Molecules as Documents." *Journal of Theoretical Biology* 8.

FIGURES

Figure 1. The steps involved in a Bayesian phylogenetic analysis of lexical data. First the lexical data is cognate coded, then the cognate sets are expressed as either a multistate or binary matrix, then the prior information for the Bayesian analysis is specified, then the stochastic model of character change is selected, then multiple MCMC searches are run, and finally the resulting posterior distribution of trees and their associated parameters are summarized.

Figure 2: An example of the parameter trace for the posterior probability of the Central Pacific language group. Each point is one sampled generation from the MCMC chain. Initially, the posterior estimate is very low (more negative) reflecting the non-optimal randomised starting points, but increases over time. After around 2.5×10^7 generations, the posterior stabilises to be around -38500. The samples before this stabilisation (in gray) are considered burn-in and are not included in the analysis.

Figure 3. A densitree and map of the Central Pacific languages constructed from the posterior distribution of trees from a Bayesian analysis of basic vocabulary. Note that the densitree reveals considerable conflicting signal that cannot be captured in single tree.

Figure 4. A comparison of the Austronesian phylogeny (Maximum Clade Credibility tree from a Bayesian analysis) from Gray et al. (2009) vs. the “known” subgrouping from Ethnologue. The Bayesian estimate recovers most of the major subgroups in the Ethnologue classification, and the mismatches mainly occur where the putative subgroups are not broadly accepted. Both trees differ enormously from lexicostatistical estimates of Austronesian relationships (see Greenhill and Gray 2012).







